

Building a Digital Twin Network

Pere Barlet
(pere.barlet@upc.edu)

Joint work with: Paul Almasan, Guillermo Bernárdez,
Miquel Ferriol, Jordi Paillissé, José Suárez-Varela, Albert Cabellos



Questions

What is a Digital Twin Network?

What can it be useful for?

How can we build one?

How far are we from building it?

How can we leverage AI/ML for this?

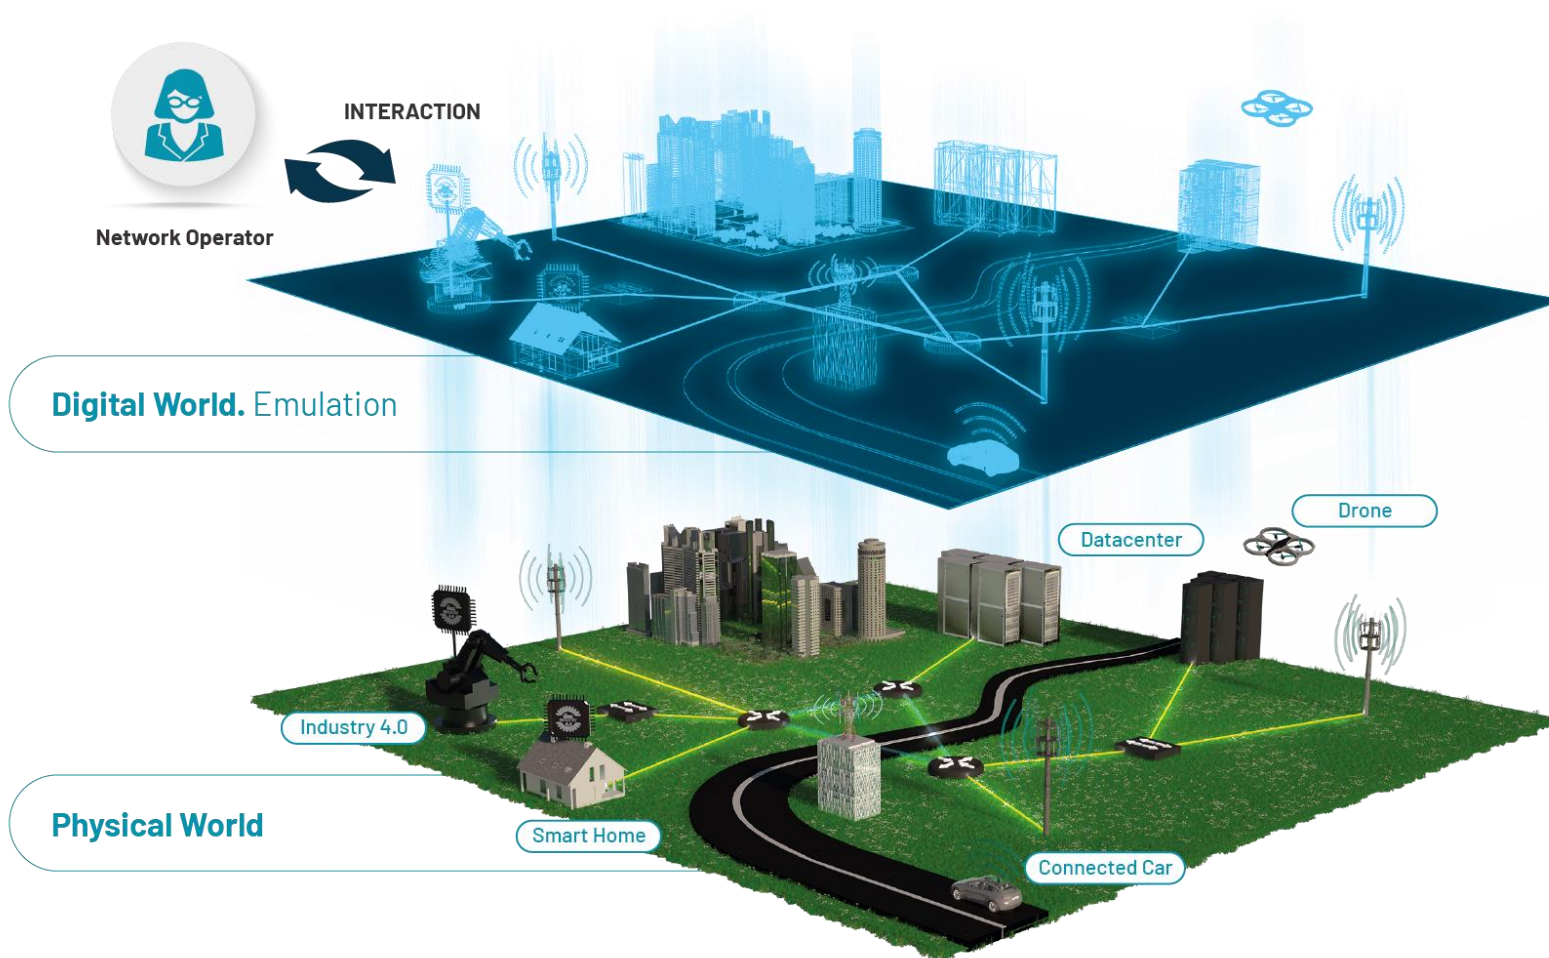
Motivation

- The Internet is a huge, complex and dynamic system
 - >20B connected devices (and growing)¹
 - 50000 GB/s of data moved daily by the Internet¹
 - 2.5 Exabytes (10^{18}) of new data created daily²
- Digital Twin paradigm
 - Virtual model of a physical object or system represented in the digital world
 - Models complex and dynamic systems
 - Examples: Aerospace, Aeronautics, Industry 4.0, etc.

¹ Cisco Annual Internet Report (2018–2023) White Paper

² <https://www.ibmbigdatahub.com/infographic/four-vs-big-data>

Digital Twin Network (DTN)



- Virtual replica of the communication network
- Model intricacies of real networks in the virtual world
- Test new configurations without compromising the actual network
- IETF, ITU-T working on it^{1,2}

¹C. Zhou et al., "Digital Twin Network: Concepts and Reference Architecture," IETF, Internet-Draft, 2022.

²ITU-T, "Digital twin network: Requirements and architecture," Recommendation ITU-T Y.3090, 2022.

DTN applications

Network planning

- Topology design
- Network dimensioning

Performance analysis

- Delay, jitter, loss
- SLA prediction
- Estimate operational bounds

What-if analysis

- Failures
- Traffic changes
- Configuration changes

Network optimization

- Traffic Engineering
- Support more SLA with current resources

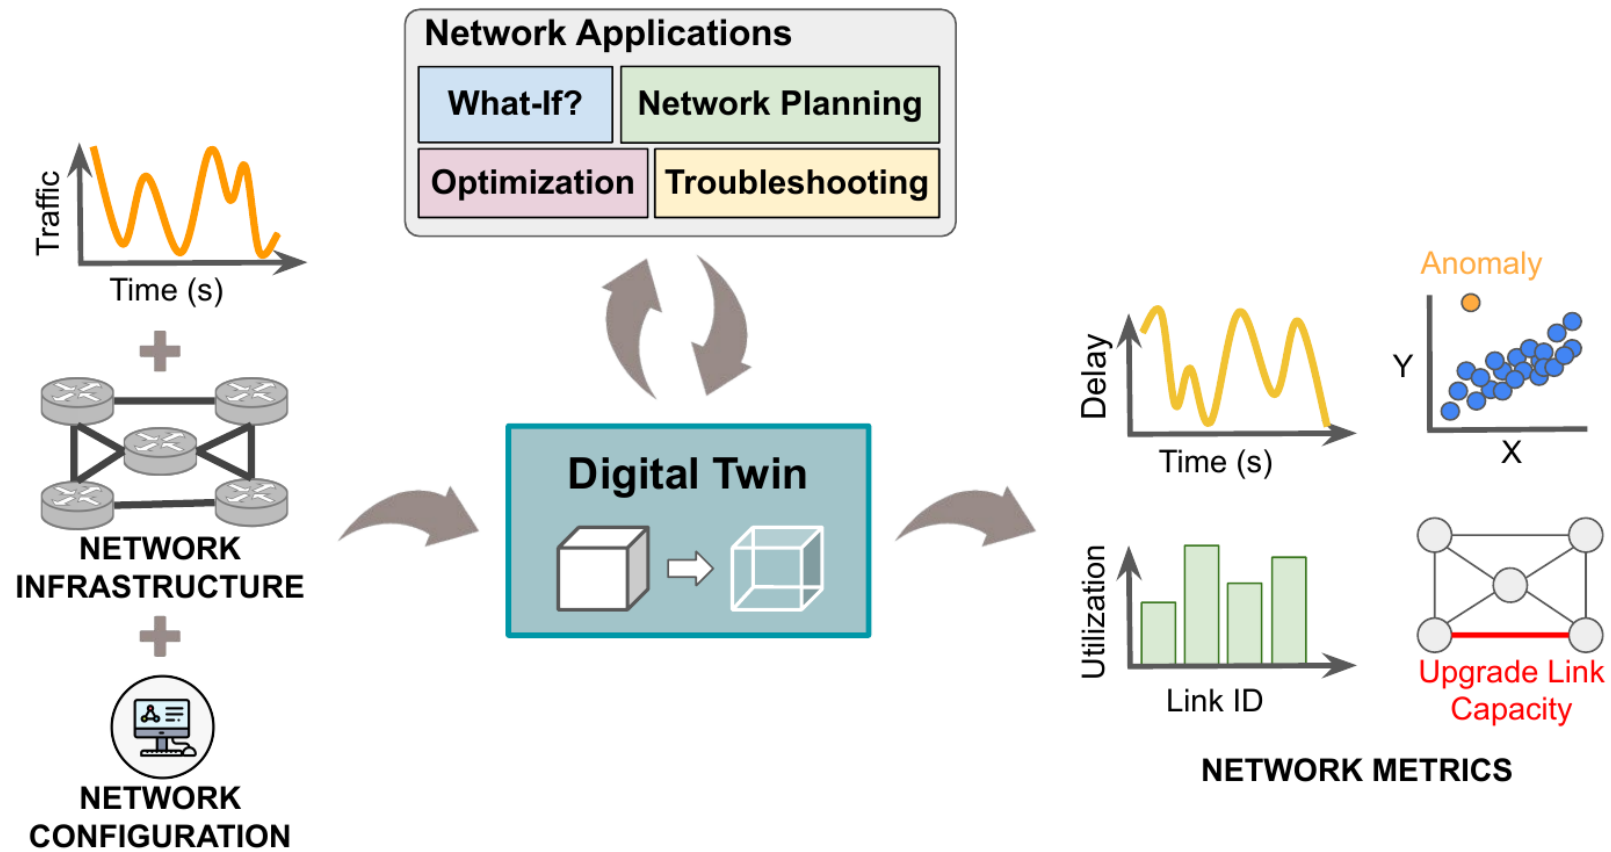
Troubleshooting

- Performance issues
- Anomaly detection

Education and training

- Certifications
- Learning

What is a DTN for us?



A DTN can be as simple or complex as required by the use case

Why does it not exist yet?

Tool	Examples	Cost	Accuracy
Analytical models	Queuing Theory, Fluid models, ...	Low	Low
Simulators	NS-3, Omnet++, Packet Tracer, ...	High	High
Emulators	GNS3, Mininet, ...	High	Low
Testbeds	Fed4FIRE, GENI, ...	High	High
Neural networks	MLP, CNN, RNN, ...	Low	???

Our research question


How can we leverage the recent advancements in **AI/ML** to build and **efficient** and **accurate** DTN?

Accurate as network
simulators

Efficient as analytical
models

Limitations of traditional NN

- Current proposals **do not generalize** to other networks
- Can they be accurate? Yes, but they require training on customer premises
 - Network instrumentation
 - Performance degradation
 - Service disruption
- Why do they not generalize?
 - Networks are essentially graphs (input)
 - Variable size input (nodes/edges)
 - Information is relational



Our vision

(1) Speed and accuracy

- Enables online optimization
- Self-driving networks, closed-loop operation

(2) Generalization

- Operation in networks never seen in training
- Enables instant deployment

(3) Scalability

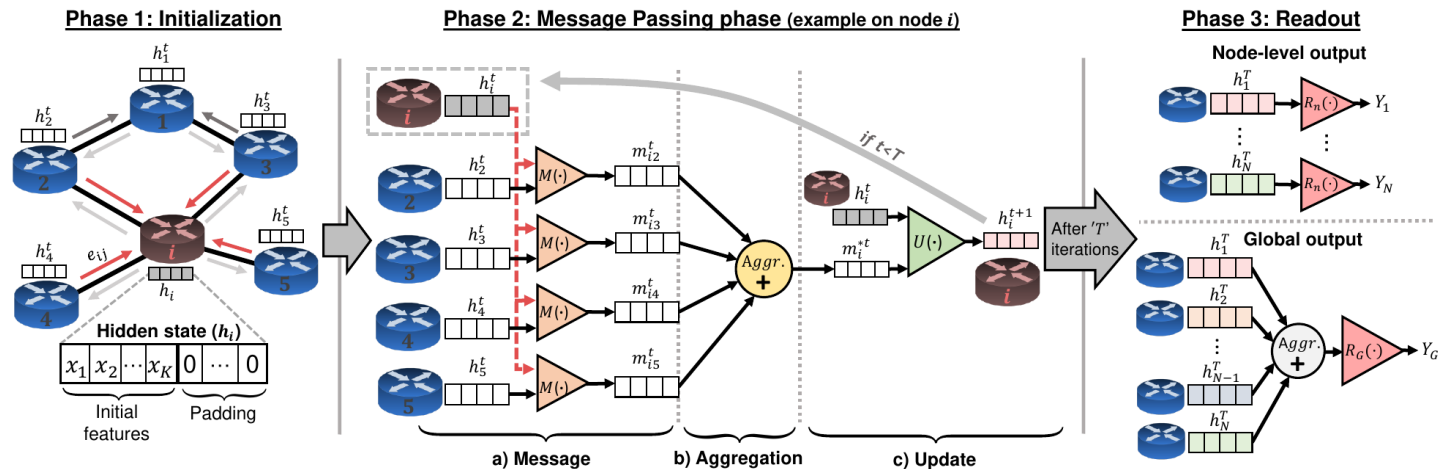
- Scale to *much* larger networks
- Enables training in a controlled testbed

(4) Accessibility

- Accessible to non-ML experts
- Enables easier adoption


Graph Neural Networks

- GNN are a *key* technology to build a DTN because:
 - Assembled dynamically based on the input graph
 - Support variable sized input graphs
 - Model relational information (e.g. graph isomorphism)
 - Generalize to different graphs

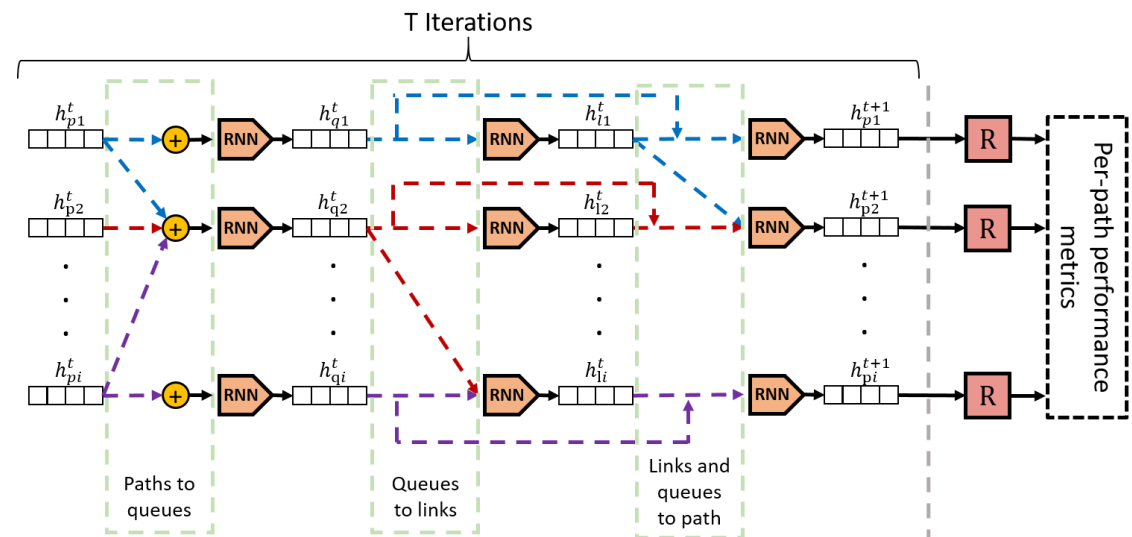


RouteNet

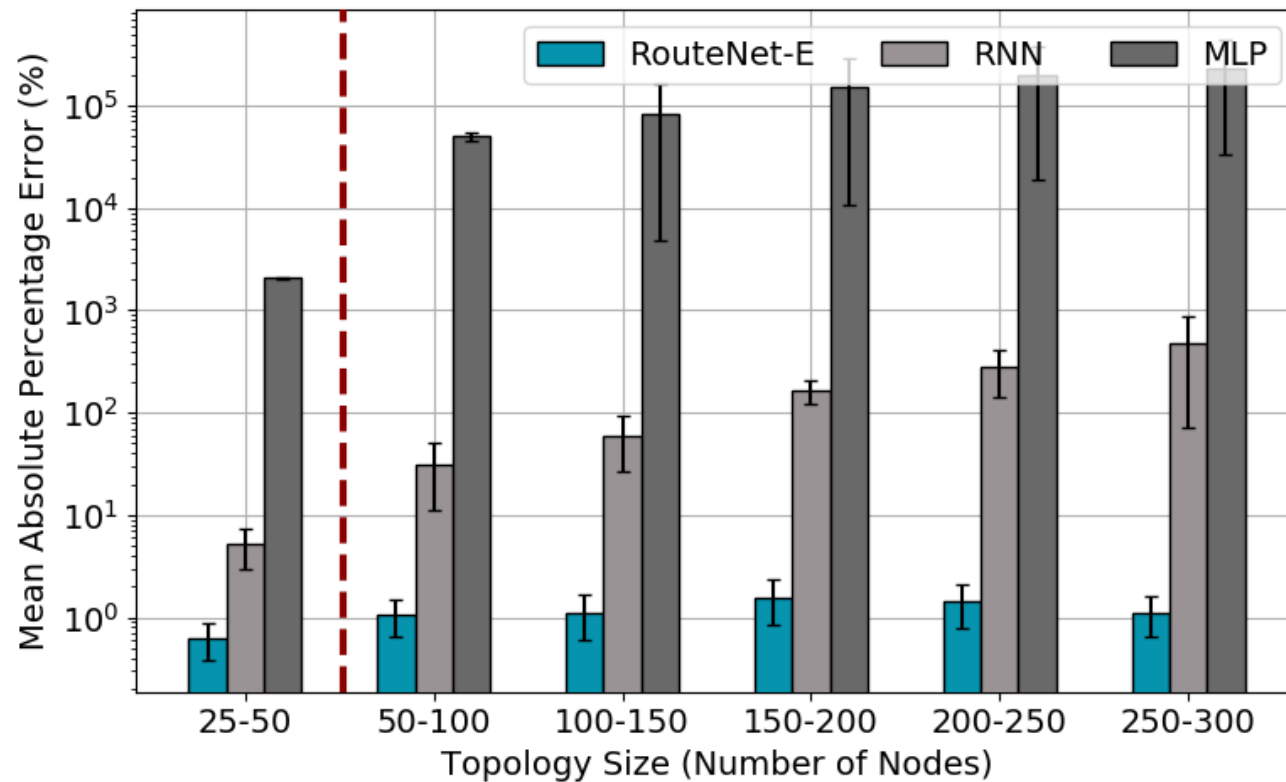
- RouteNet is the reference GNN for network modeling
 - It can operate in networks never seen in training
 - Networks several orders of magnitude larger
 - Cost equivalent to analytical models
 - Accuracy equivalent to simulation



Currently at version E
RouteNet-E (RouteNet-Erlang)

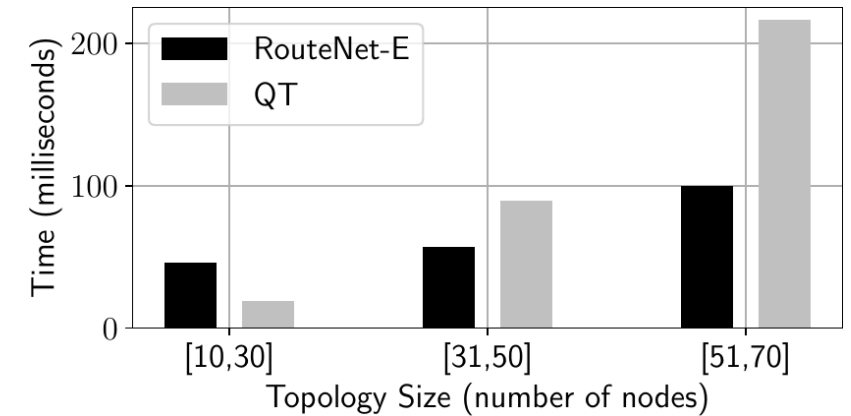


RouteNet-E: Performance

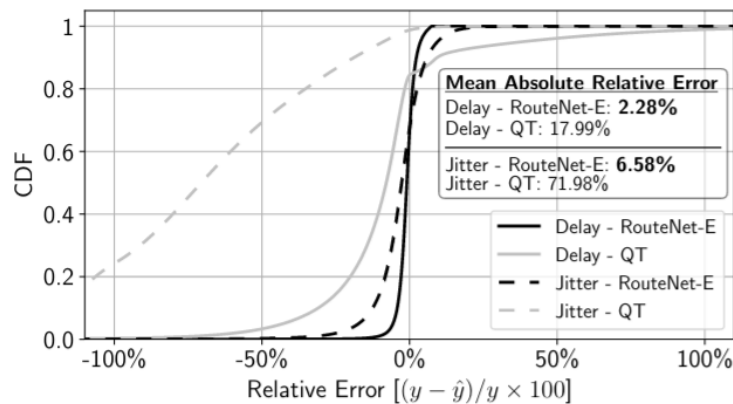


Experiment setup:


- Training: 100K samples, 25-50 nodes simultaneously
- Test: 500 samples 50-300 nodes uniformly




RouteNet-E: Performance




(a) Poisson




(b) Constant bitrate




(c) On-Off



(d) Autocorrelated exponentials



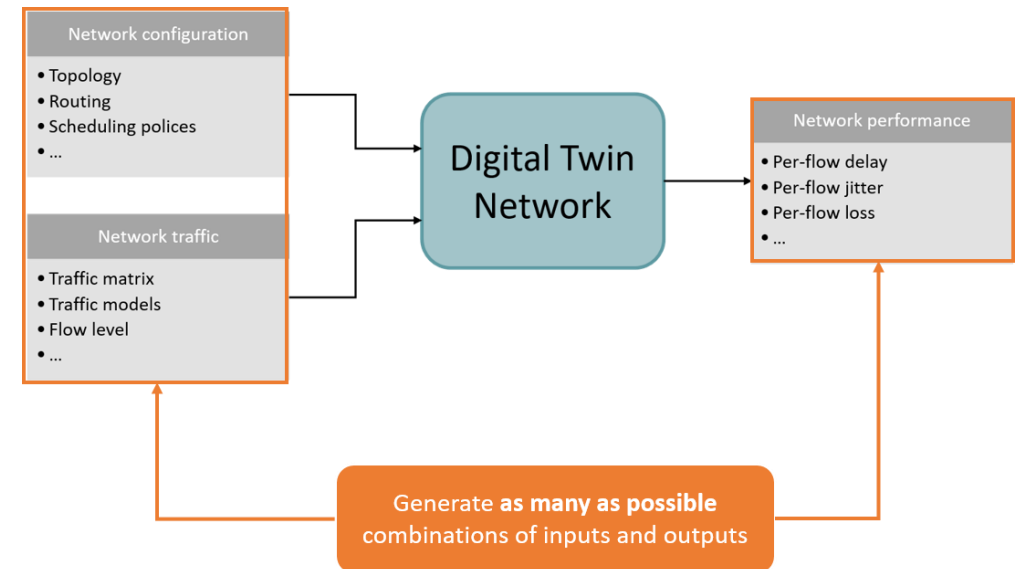
(e) Modulated exponentials



(f) All traffic models multiplexed

How to train a DTN?

- Previous ML approaches proposed **online training** (or transfer learning)
 - Not practical from a product point of view
 - Training models is a huge effort, requires a large data set and instrumentation
 - Can lead to degraded performance or service disruption
- RouteNet enables **offline training** in a controlled testbed or simulator
 - Can cover all possible situations
 - Models can be validated before deployment
 - Can provide safety bounds



Conclusions

Tool	Examples	Cost	Accuracy
Analytical models	Queuing Theory, Fluid models, ...	Low	Low
Simulators	NS-3, Omnet++, Packet Tracer, ...	High	High
Emulators	GNS3, Mininet, ...	High	Low
Testbeds	Fed4FIRE, GENI, ...	High	High
<i>Traditional</i> Neural networks	MLP, CNN, RNN, ...	Low	???

Conclusions

Tool	Examples	Cost	Accuracy
Analytical models	Queuing Theory, Fluid models, ...	Low	Low
Simulators	NS-3, Omnet++, Packet Tracer, ...	High	High
Emulators	GNS3, Mininet, ...	High	Low
Testbeds	Fed4FIRE, GENI, ...	High	High
<i>Traditional</i> Neural networks	MLP, CNN, RNN, ...	Low	Low

Conclusions

Tool	Examples	Cost	Accuracy
Analytical models	Queuing Theory, Fluid models, ...	Low	Low
Simulators	NS-3, Omnet++, Packet Tracer, ...	High	High
Emulators	GNS3, Mininet, ...	High	Low
Testbeds	Fed4FIRE, GENI, ...	High	High
<i>Traditional</i> Neural networks	MLP, CNN, RNN, ...	Low	Low
Graph Neural Networks	RouteNet-E	Low	High



Learn: All papers are free online
<https://github.com/BNN-UPC/Papers/wiki>



Play: Code and Datasets open-source
<https://bnn.upc.edu>




Code: IGNNITION framework
<https://ignnition.net>



Challenge yourself: GNN Challenge
<https://bnn.upc.edu/challenge>

Backup slides


What is a DTN for us?



A DTN can be as simple or complex as needed

Limitations of traditional NN

- Why do they not generalize?
 - Networks are essentially graphs (input data)
 - Graphs are of variable size (links and nodes)
 - Information is relational (as opposed to Euclidian or sequential)
 - Modeling networks with traditional NN is **very hard!**



RouteNet-E: Algorithm

Algorithm 1 Internal architecture of RouteNet-E

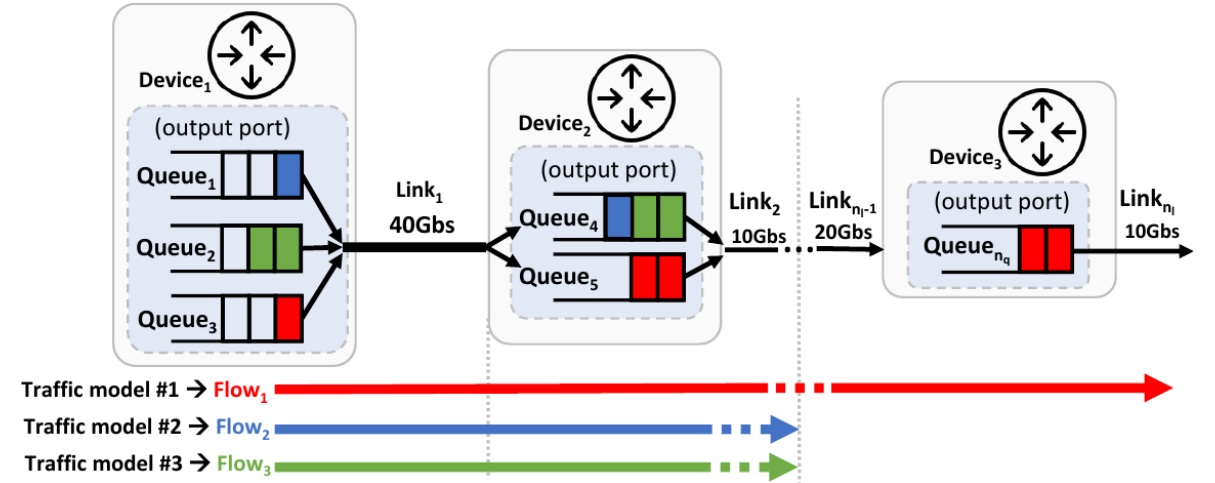
Input: F, Q, L, x_f, x_q, x_l

Output: $h_q^T, h_l^T, h_f^T, \hat{y}_f, \hat{y}_q, \hat{y}_l$

```



1: for each  $l \in L$  do  $h_l^0 \leftarrow [x_l, 0...0]$ 
2: for each  $q \in Q$  do  $h_q^0 \leftarrow [x_q, 0...0]$ 
3: for each  $f \in F$  do  $h_f^0 \leftarrow [x_f, 0...0]$ 
4: for  $t = 0$  to  $T-1$  do
5:   for each  $f \in F$  do                                     ▷ Message Passing Phase
6:     for each  $(q, l) \in f$  do                               ▷ Message Passing on Flows
7:        $h_f^t \leftarrow FRNN(h_f^t, [h_q^t, h_l^t])$            ▷ Flow: Aggr. and Update
8:        $\tilde{m}_{f,q}^{t+1} \leftarrow h_f^t$                        ▷ Flow: Message Generation
9:        $h_f^{t+1} \leftarrow h_f^t$ 
10:    for each  $q \in Q$  do                                     ▷ Message Passing on Queues
11:       $M_q^{t+1} \leftarrow \sum_{f \in Q_f(q)} \tilde{m}_{f,q}^{t+1}$        ▷ Queue: Aggregation
12:       $h_q^{t+1} \leftarrow U_q(h_q^t, M_q^{t+1})$              ▷ Queue: Update
13:       $\tilde{m}_q^{t+1} \leftarrow h_q^{t+1}$                        ▷ Queue: Message Generation
14:    for each  $l \in L$  do                                     ▷ Message Passing on Links
15:      for each  $q \in L_q(l)$  do
16:         $h_l^t \leftarrow LRNN(h_l^t, \tilde{m}_q^{t+1})$            ▷ Link: Aggr. and Update
17:         $h_l^{t+1} \leftarrow h_l^t$ 
18:   $\hat{y}_f \leftarrow R_f(h_f^T)$                                ▷ Readout phase
19:   $\hat{y}_q \leftarrow R_q(h_q^T)$ 

```



RouteNet-E: Scalability

- Out-of-distribution
 - Larger paths
 - Larger link capacities
 - Output distributions (delay)

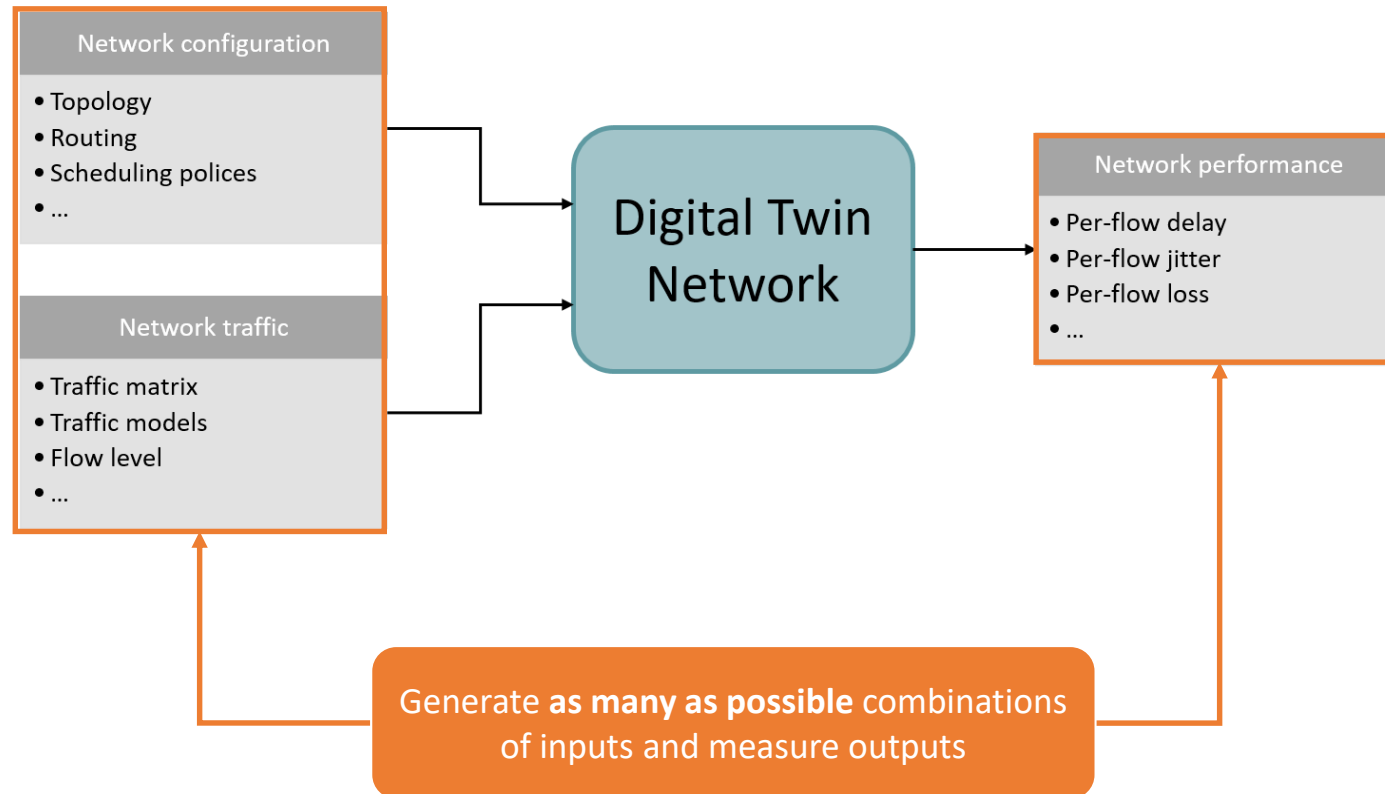


$$\text{Delay}_{L_1} = \text{Avg_utilization}_{Q_1} * \text{Size}_{Q_1} / \text{Cap}_{L_1}$$

$$\text{Delay}_{\text{Path}} = \sum_{k=1}^{N \text{ links}} \text{Delay}_{L_k}$$

Our solution


- Small-scale testbed at **vendor** premises (before deployment)
- Leverage GNN generalization/scalability features




Data sets need to contain all possible scenarios, including those that result in poor performance or service disruption !!

Multi-Agent Distributed Optimization


- Agent that operates networks autonomously
- Exchanges messages (protocol) with other agents
- Uses a GNN to learn message contents
- Optimizes network performance using Multi-Agent RL
 - Example: Traffic Optimization, Congestion, etc.
- Scales well as it naturally (GNN) distributes load among devices
- **Supports offline learning, can operate in unseen networks**



MARL+GNN: Optimization



(a) TopologyZoo Uniform



(b) TopologyZoo Gravity

MARL (Proximal Policy Optimization) + GNN
100 TM (uniform / gravity)
75 topologies from Topology Zoo
Trained with 2 topologies (NSFnet, GEANT2)